

Next-Word-Prediction

Step 1: Import libraries

```
In [1]: ▶ import tensorflow as tf
        from tensorflow.keras.preprocessing.text import Tokenizer
        from tensorflow.keras.layers import Embedding, LSTM, Dense
        from tensorflow.keras.models import Sequential
        from tensorflow.keras.utils import to_categorical
        from tensorflow.keras.optimizers import Adam
        import pickle
        import numpy as np
        import os
        !pip install pydot
        import matplotlib.pyplot as plt
```

Requirement already satisfied: pydot in c:\users\snigdha\anaconda3\lib\site-packages (1.4.2)

Requirement already satisfied: pyparsing>=2.1.4 in c:\users\snigdha\anaconda3\lib\site-packages (from p
ydot) (3.0.9)

Step 2: Load your file

```
In [2]: ▶ path = "C:\\Users\\Snigdha\\OneDrive\\Desktop\\New folder\\pg145.txt"
```

Step 3: Open and pre-process the data

```
In [3]: file = open("C:\\Users\\Snigdha\\OneDrive\\Desktop\\New folder\\pg145.txt", "r", encoding = "utf8")

# store file in list
lines = []
for i in file:
    lines.append(i)

# Convert list to string
data = ""
for i in lines:
    data = ' '.join(lines)

#replace unnecessary stuff with space
data = data.replace('\n', ' ').replace('\r', ' ').replace('\ufeff', ' ').replace('"', ' ').replace("'", ' ') #r

#remove unnecessary spaces
data = data.split()
data = ' '.join(data)
data[:500]
```

Out[3]: 'The Project Gutenberg eBook of Middlemarch, by George Eliot This eBook is for the use of anyone anywhere in the United States and most other parts of the world at no cost and with almost no restrictions whatsoever. You may copy it, give it away or re-use it under the terms of the Project Gutenberg License included with this eBook or online at www.gutenberg.org. If you are not located in the United States, you will have to check the laws of the country where you are located before using this eBook'

```
In [5]: len(data)
```

Out[5]: 1783493

Step 4: Apply tokenization and some other changes

```
In [6]: ▶ tokenizer = Tokenizer()
tokenizer.fit_on_texts([data])

# saving the tokenizer for predict function
pickle.dump(tokenizer, open('token.pkl', 'wb'))

sequence_data = tokenizer.texts_to_sequences([data])[0]
sequence_data[:15]
```

```
Out[6]: [1, 366, 358, 1985, 3, 158, 32, 2247, 5558, 35, 1985, 24, 20, 1, 258]
```

```
In [7]: ▶ len(sequence_data)
```

```
Out[7]: 321821
```

```
In [8]: ▶ vocab_size = len(tokenizer.word_index) + 1
print(vocab_size)
```

```
17871
```

```
In [9]: ► sequences = []

for i in range(3, len(sequence_data)):
    words = sequence_data[i-3:i+1]
    sequences.append(words)

print("The Length of sequences are: ", len(sequences))
sequences = np.array(sequences)
sequences[:10]
```

The Length of sequences are: 321818

```
Out[9]: array([[ 1, 366, 358, 1985],
 [366, 358, 1985,  3],
 [358, 1985,  3, 158],
 [1985,  3, 158, 32],
 [ 3, 158, 32, 2247],
 [158, 32, 2247, 5558],
 [ 32, 2247, 5558, 35],
 [2247, 5558, 35, 1985],
 [5558, 35, 1985, 24],
 [ 35, 1985, 24, 20]])
```

```
In [10]: ► X = []
y = []

for i in sequences:
    X.append(i[0:3])
    y.append(i[3])

X = np.array(X)
y = np.array(y)
```

```
In [11]: ▶ print("Data: ", X[:10])
          print("Response: ", y[:10])
```

```
Data: [[ 1 366 358]
 [ 366 358 1985]
 [ 358 1985 3]
 [1985 3 158]
 [ 3 158 32]
 [ 158 32 2247]
 [ 32 2247 5558]
 [2247 5558 35]
 [5558 35 1985]
 [ 35 1985 24]]
Response: [1985 3 158 32 2247 5558 35 1985 24 20]
```

```
In [12]: ▶ y = to_categorical(y, num_classes=vocab_size)
          y[:5]
```

```
Out[12]: array([[0., 0., 0., ..., 0., 0., 0.],
 [0., 0., 0., ..., 0., 0., 0.],
 [0., 0., 0., ..., 0., 0., 0.],
 [0., 0., 0., ..., 0., 0., 0.],
 [0., 0., 0., ..., 0., 0., 0.]], dtype=float32)
```

Step 5: Creating the model

```
In [13]: ▶ model = Sequential()
          model.add(Embedding(vocab_size, 10, input_length=3))
          model.add(LSTM(1000, return_sequences=True))
          model.add(LSTM(1000))
          model.add(Dense(1000, activation="relu"))
          model.add(Dense(vocab_size, activation="softmax"))
```

In [15]: `model.summary()`

Model: "sequential"

Layer (type)	Output Shape	Param #
embedding (Embedding)	(None, 3, 10)	178710
lstm (LSTM)	(None, 3, 1000)	4044000
lstm_1 (LSTM)	(None, 1000)	8004000
dense (Dense)	(None, 1000)	1001000
dense_1 (Dense)	(None, 17871)	17888871


=====
Total params: 31,116,581
Trainable params: 31,116,581
Non-trainable params: 0
=====

Step 6: Plot the model

```
In [16]: from tensorflow import keras
from keras.utils.vis_utils import plot_model

keras.utils.plot_model(model, to_file='plot.png', show_layer_names=True)
```

You must install pydot (``pip install pydot``) and install graphviz (see instructions at <https://graphviz.gitlab.io/download/>) (<https://graphviz.gitlab.io/download/>) for plot_model to work.



Step 7: Train the model

```
In [17]: ▶ from tensorflow.keras.callbacks import ModelCheckpoint
from tensorflow.keras import Model

checkpoint = ModelCheckpoint("next_words.h5", monitor='loss', verbose=1, save_best_only=True)
```

```
In [18]: ▶ model.compile(loss="categorical_crossentropy", optimizer=Adam(learning_rate=0.001))
```

```
In [19]: ▶ history = model.fit(X, y, epochs=3, batch_size=64, callbacks=[checkpoint])
```

```
Epoch 1/3
5029/5029 [=====] - ETA: 0s - loss: 6.3798
Epoch 1: loss improved from inf to 6.37982, saving model to next_words.h5
5029/5029 [=====] - 666s 131ms/step - loss: 6.3798
Epoch 2/3
5029/5029 [=====] - ETA: 0s - loss: 5.7952
Epoch 2: loss improved from 6.37982 to 5.79525, saving model to next_words.h5
5029/5029 [=====] - 704s 140ms/step - loss: 5.7952
Epoch 3/3
5029/5029 [=====] - ETA: 0s - loss: 5.5164
Epoch 3: loss improved from 5.79525 to 5.51639, saving model to next_words.h5
5029/5029 [=====] - 680s 135ms/step - loss: 5.5164
```

Step-8:Lets Predict


```
In [21]: ▶ from tensorflow.keras.models import load_model
import numpy as np
import pickle

# Load the model and tokenizer
pickle.dump(tokenizer, open('token.pkl', 'wb'))
model = load_model('next_words.h5')
tokenizer = pickle.load(open('token.pkl', 'rb'))

def Predict_Next_Words(model, tokenizer, text):

    sequence = tokenizer.texts_to_sequences([text])
    sequence = np.array(sequence)
    preds = np.argmax(model.predict(sequence))
    predicted_word = ""

    for key, value in tokenizer.word_index.items():
        if value == preds:
            predicted_word = key
            break

    print(predicted_word)
    return predicted_word
```

```
In [22]: ► while(True):
          text = input("Enter your line: ".lower())

          if text == "0":
              print("Execution completed.....")
              break

          else:
              try:
                  text = text.split(" ")
                  text = text[-3:]
                  print(text)

                  Predict_Next_Words(model, tokenizer, text)

              except Exception as e:
                  print("Error occurred: ",e)
                  continue
```

```
enter your line: If you are not located in the United States, you will have to check the laws of the
['laws', 'of', 'the']
1/1 [=====] - 2s 2s/step
world
enter your line: and deed in noble agreement; but after all, to common eyes their struggles
['their', 'struggles', '']
1/1 [=====] - 1s 1s/step
of
enter your line: Meanwhile the indefiniteness remains, and the limits of variation are really much wide
r than
['wider', 'than', '']
1/1 [=====] - 0s 31ms/step
the
enter your line: The Project Gutenberg eBook of
['eBook', 'of', '']
1/1 [=====] - 0s 32ms/step
the
enter your line: This eBook is for the use of anyone anywhere
['of', 'anyone', 'anywhere']
1/1 [=====] - 0s 47ms/step
and
enter your line: 0
Execution completed.....
```

In []: ▶